



# C++ Micro-Problems

**Davide Di Gennaro**

Software Engineer



# C++ Micro-Problems

- Herb Sutter (2013) “My Favorite C++ 10-Liner”  
<https://channel9.msdn.com/Events/GoingNative/2013/My-Favorite-Cpp-10-Liner>
  - Indeed some (\*) problems can be solved by 10 lines of C++
  - We will show you three of them
- 
- (\*) Not *all* problems unfortunately, but we are working on it

# SLIDWARE

- All the problems described here admit endless variations
- Give a working solution to a problem, but not necessarily complete
- Some important details/optimizations have been intentionally omitted

# Global Parameters

Problem:

- There is a global variable that rarely changes
- We need to read it often, from multiple threads
- assume it's not a basic type, otherwise `std::atomic` suffices

# Global Parameters

## Solution #1

```
mutex m;
string theName;

void SetTheName(const string& s)
{
    lock_guard g(m);
    theName = s;
}

string GetTheName()
{
    lock_guard g(m);
    return theName;
}
```

# Global Parameters

## Solution #1

- Simple but not super-efficient
  - Locks everytime
  - Copies the string everytime

Ok, Let's try another approach...

# Global Parameters

## Solution #2

```
#include <list>
list<string> theHistory{ "" };

void SetTheName(const string& s)
{
    theHistory.push_front(s);
}

const string& GetTheName()
{
    return theHistory.front();
}
```

# Global Parameters

## Solution #2

- Interesting approach...
  - Not minimally thread safe!
  - We removed the copy but we now have no way of deleting old values

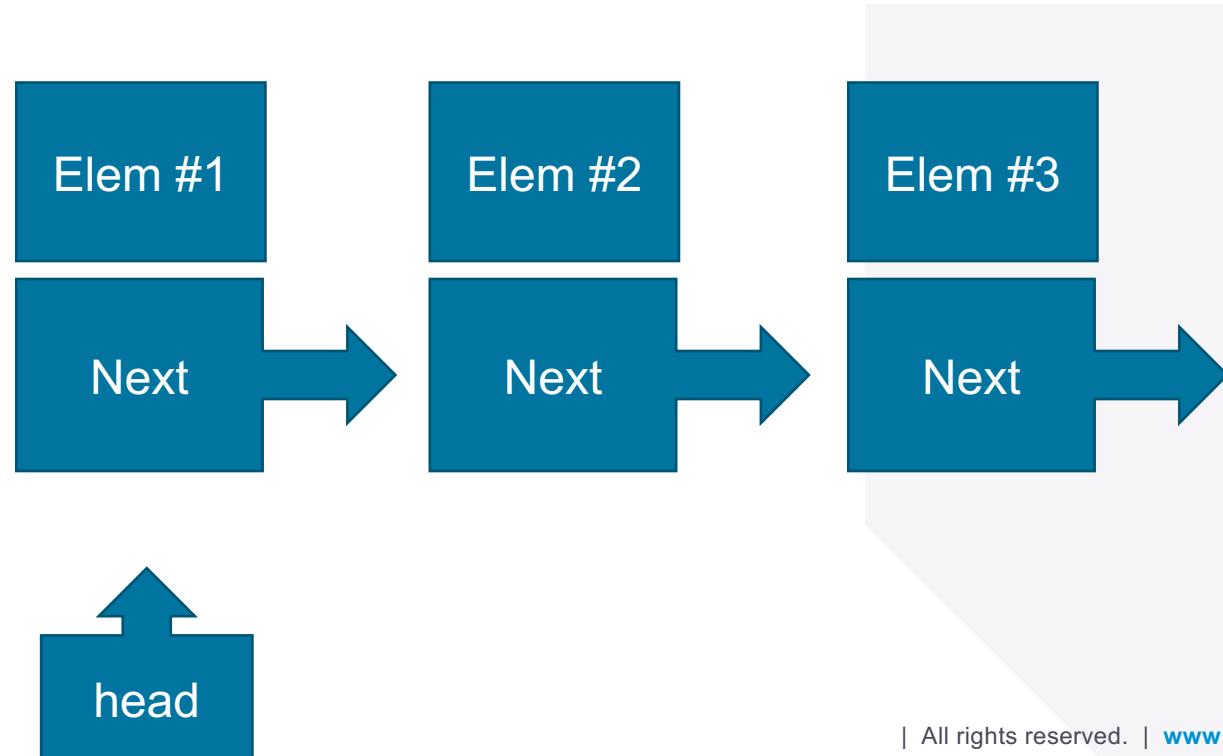
# Global Parameters

## Solution #2

- Interesting approach... but we can do better
  - Not minimally thread safe!
  - We removed the copy but we now have no way of deleting old values
- Under these assumptions we can use an **atomic** linked list

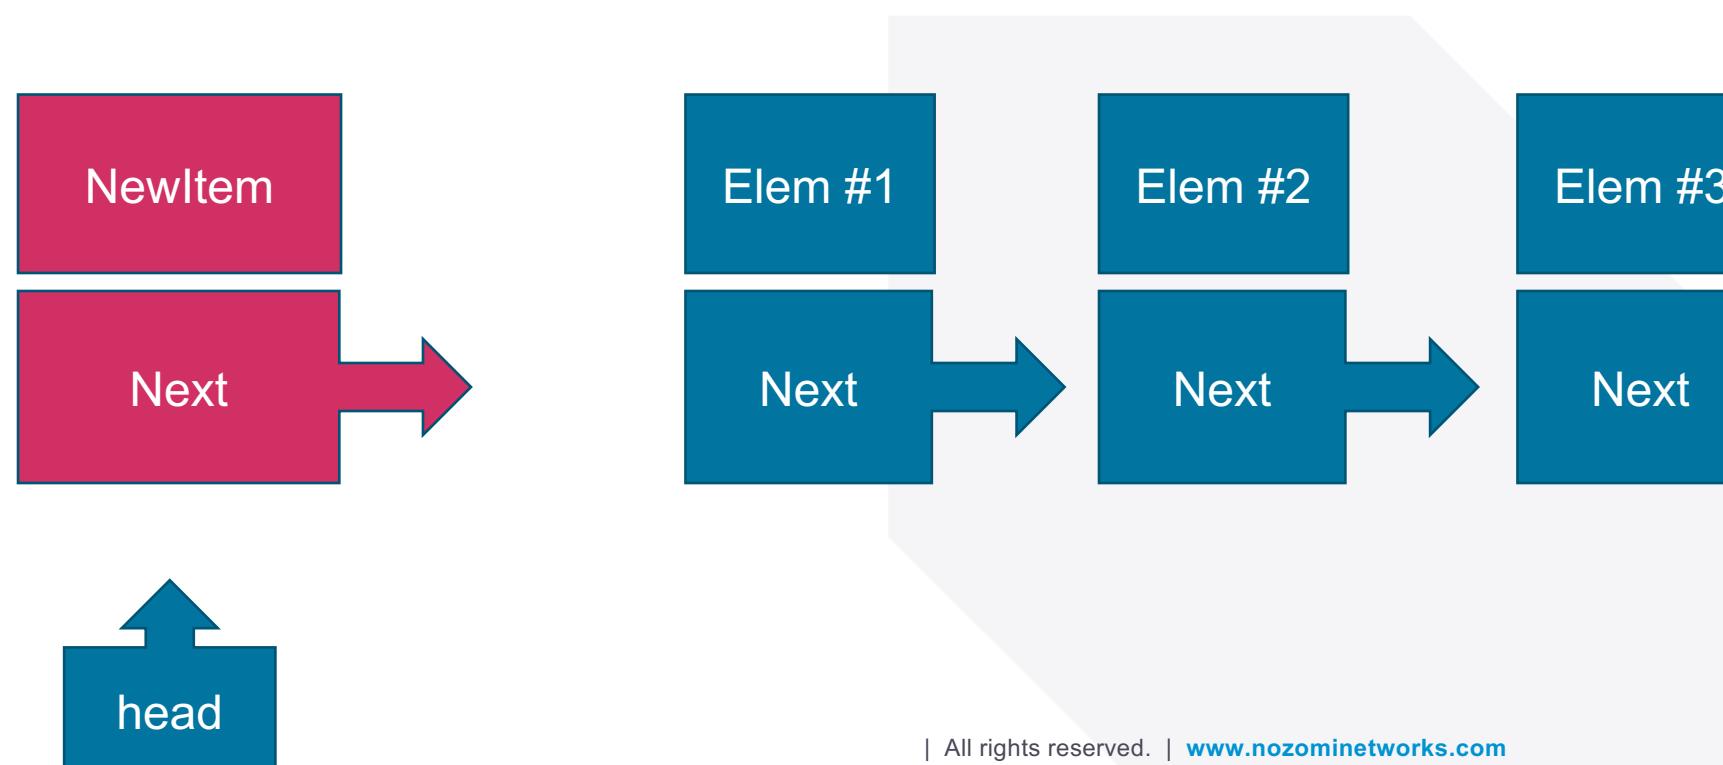
This is a feature  
not a bug

# Global Parameters



# Global Parameters

- To insert at front, **atomically-swap head and head→next**



# Global Parameters

```
#include "atomic_list.h"
AtomicList<string> theHistory;

void SetTheName(const string& s)
{
    theHistory.PushFront(s);
}

const string& GetTheName()
{
    return theHistory.Front();
}
```

# Global Parameters

```
template <typename T>
class AtomicList
{
    struct Node
    {
        T value;
        std::atomic<Node*> next {nullptr};
    };

    std::atomic<Node*> head_ {nullptr};

public:

    template <typename U>
    void PushFront(U&& x);

    const T& Front();
};
```

# Global Parameters

atomic\_list.h

```
template<typename U>
class AtomicList
{
    Node* const n = new Node {std::forward<U>(x), nullptr};
    struct Node* h = head_.load();
    do
    {
        T value;
        n->next = h;
        std::atomic<Node*>.next {nullptr};
    } while (!head_.compare_exchange_weak(h, n));
};

std::atomic<Node*> head_ {nullptr};
{
    static const T empty {};
public: if (auto s = head_.load())
{
    template<typename U>
    void PushFront(U&& x);
    else
    {
        const T& return empty;
    }
}
};
```

**YO DAWG, HEARD  
YOU LIKE ATOMICS**



**SO I PUT AN ATO...**

# Global Dictionary

Problem:

- we have a large number of strings but the set of different values is small

Example:

```
class Node
{
    std::string vendor_;
```

- Remember that an empty string uses ~ 24 bytes
- If 95% of nodes have vendor "", the set of vendors qualifies as small

# Global Dictionary

- String Interning
- A technique made popular by Java
- String instances are added to a global dictionary and a pointer to the dictionary entry is used instead
- Bonus: can test equality very fast

# Global Dictionary

- String Interning
- A technique made popular by Java
- String instances are added to a global dictionary and a pointer to the dictionary entry is used instead
- Bonus: can test equality very fast

```
template <typename T>
struct GlobalDictionary
{
    using dictionary_t = std::unordered_set<std::string>;
    static const std::string* Intern(const std::string& s)
    {
        static dictionary_t words;
        static std::mutex m;
        std::lock_guard g(m);
        return &(*words.insert(s).first);
    }

    static const std::string* Empty()
    {
        static const std::string e;
        return &e;
    }
};
```

Pass a string

exchange it for  
a pointer into dictionary

# Global Dictionary

GlobalString.h

- String Interning
- A technique made popular by Java
- String instances are added to a global dictionary and a pointer to the dictionary entry is used instead
- Bonus: can test equality very fast
- The pointer is wrapped in a pseudo-string object

```
class GlobalString
{
    const std::string* str_;
public:
    GlobalString()
        : str_(GlobalDictionary<void>::Empty())
    {
    }

    GlobalString(const std::string& s)
        : str_(GlobalDictionary<void>::Intern(s))
    {
    }

    // ...
}
```

# Global Dictionary

GlobalString.h

- String Interning
- A technique made popular by Java
- String instances are wrapped in a GlobalDictionary and a pointer to the dictionary entry is used instead of the string object
- Bonus: can test equality directly
  - return ptr\_ == that.ptr\_ || \*ptr\_ == \*that.ptr\_;
- The pointer is wrapped in a GlobalString

```
class GlobalString
{
    const std::string* str_;

public:
    GlobalString()
        : str_(GlobalDictionary<void>::Empty())
    {
    }

    bool operator==(const GlobalString& that) const
    {
        return ptr_ == that.ptr_ || *ptr_ == *that.ptr_;
    }

    GlobalString(const std::string& s)
        : str_(GlobalDictionary<void>::Intern(s))
    {
    }

    // ...
}
```

# Global Dictionary

- String Interning
- A technique made popular by Java
- String instances are added to a global dictionary and a pointer to the dictionary entry is used instead
- Bonus: can test equality very fast
- The pointer is wrapped in a pseudo-string object
- Bonus: use thread\_local storage, but global life!

```
template <typename T>
struct GlobalDictionary
{
    using dictionary_t = std::unordered_set<std::string>;
```

```
    template <typename X>
    static X& CreateAnotherFrom(std::list<X>& c)
    {
        static std::mutex m;
        std::lock_guard g(m);
        return c.emplace_back();
    }
```

```
    static const std::string* Intern(const std::string& s)
    {
        static std::list<dictionary_t> global;
```

```
        static thread_local dictionary_t& words = CreateAnotherFrom(global);
        return &(*words.insert(s).first);
    }
```

```
    static const std::string* Empty()
    {
        static const std::string e;
        return &e;
    }
};
```



borrow local storage  
from a global container



**WHICH THREAD, PRECISELY?**

# Global Object Tracker

Problem:

- A static variable is initialized when the process starts
- C++ unit tests run in a single process that executes all test functions in some order
- Sometimes a reshuffle makes test fail

# Global Object Tracker

Example:

nTasks

nTasks++

Test.cpp

```
//class AsyncTaskQueue
//{
//    static int nTasksProcessed = 0;
//    ...

class AsyncTaskQueueTest : public ::testing::Test
{
};

TEST_F(AsyncTaskQueueTest, Test1)
{
    AsyncTaskQueue q;
    q.Enqueue( [] { DoSomething(); } );
}

TEST_F(AsyncTaskQueueTest, Test2)
{
    ASSERT_EQ(0, AsyncTaskQueue::nTasksProcessed);
}
```

0 or 1

depending on the test order

# Global Object Tracker

Idea:

- With some metaprogramming tricks:
  - Intercept variable initialization
  - Create a copy of the initializer in a global list
  - When needed, ask the global list to “replay” all the initialitiazions
- Other approaches are possible, but less desirable
  - Balance reasonable coverage/ease of customization
  - Alternative: change the type of the global variable (a wrapper)



G

```
#include <string>

std::string s = global::initialize(INITIALIZE_TOKEN, &s, "abc");
std::string t = global::initialize_multi(INITIALIZE_TOKEN, &t, 10, 'x');

int a = global::initialize(INITIALIZE_TOKEN, &a, 7);
int b = global::initialize(INITIALIZE_TOKEN, &b, 8);
int c = global::initialize(INITIALIZE_TOKEN, &c);
```

int main()

```
{
```

a = 1234;

b = 789;

c = -9999;

s = "blah";

t.clear();

global::reset\_all();

```
}
```

IN

# Global Object Tracker

We will make use of two unusual language features:

1. The name of a variable is available on the right side of the declaration  
`int n = f(&n);`
2. Multiple instances of the same lambda have different types

# Global Object Tracker

```
#include <tuple>
#include <functional>

class global
{
    struct initializer_t
    {
        std::function<void()> reset;
        initializer_t* next = nullptr;
    };

    template <typename L>
    static initializer_t& get()
    {
        static initializer_t v;
        return v;
    }

    static initializer_t& first()
    {
        return get<void>();
    }

    static void push_front(initializer_t& v)
    {
        initializer_t& head = first();
        auto n = head.next;
        head.next = &v;
        v.next = n;
    }

public:

    template <typename L, typename X, typename... T>
    static X initialize(L, X* x, T... val)
    {
        initializer_t& v = get<L>();
        v.reset = [x, t = std::make_tuple(val...)] { *x = std::make_from_tuple<X>(t); };
        push_front(v);
        return X(val...);
    }

    static void reset_all()
    {
        for (auto v = first().next; v; v = v->next)
            v->reset();
    }
};

#define INITIALIZE_TOKEN []{}
```

Glo

```
#include <tuple>
#include <functional>

class global
{
    struct initializer_t
    {
        std::function<void()> reset;
        initializer_t* next = nullptr;
    };

    //...

public:
    static void reset()
    {
        for (auto v = head->next; v; v = v->next)
            v->reset();
    }
};
```

linked list  
of functions

# Global Object Tracker

```
//...

template <typename L>
static initializer_t& get()
{
    static initializer_t v;
    return v;
}

static initializer_t& first()
{
    return get<void>();
}

static void add_initializer(initializer_t& v)
{
    initializer_t* head = first();
    auto n = new initializer_t;
    head.next = n;
    n.next = v.next;
    v.next = n;
}
```

list nodes are static variables  
associated to types

# Global Object Tracker

```
public:  
  
    template <typename L, typename X, typename... T>  
    static X initialize(L, X* x, T... val)  
    {  
        initializer_t& v = get<L>();  
        v.reset = [x, t = std::make_tuple(val...)] { *x = std::make_from_tuple<X>(*t); };  
        push_front(v);  
        return X(val...);  
    }  
  
};  
  
#define INITIALIZE_TOKEN []{}
```

lambda type L as tag

# Global Object Tracker – What I didn't tell you

- This implementation is just the shortest
  - Non-intrusive
  - Good compromise between compactness/specialization/customization
- How to improve:
  - Change `push_front` into `push_back`
  - Write another `initialize` for the default constructor
    - Instead of `*x = X{}`, try calling `x->clear()`
  - Accept a lambda as initializer:
    - `std::string s = global::initialize([] { return "abc"; }, &s);`

# Global Object Tracker – What I didn't tell you

global.h

```
template <typename X>
static decltype(std::declval<X>().clear()) clear(X& x, X*)
```

```
{
```

```
    return x.clear();
```

- This implementation is just the shortest

- Non-intrusive

```
template <typename X>
```
- Good ~~static void clear(X& x, void\*)~~ compactness/specialization/customization

```
{
```

```
    x = X{};
```

- How to improve:

- Change ~~push\_front~~ into ~~push\_back~~

```
static void reset(X& x)
```
- Write another ~~initialize~~ for the default constructor
  - Instead of ~~clear(x, &x);~~ {}, try calling x->clear()

- Accept a lambda as initializer:
  - std::string s = global::initialize([] { return "abc"; }, &s);

```
template <typename X>::initialize([] { return "abc"; }, &s);
```

```
static void reset(std::atomic<X>& x)
```

```
{
```

```
    x = X{};
```

# Global

- This im

- Non
- Goo

- How to

- Cha
- Writ
- Acces

```
template <typename L, typename X, typename... T>
static X initialize(L lambda, X* x, T... val)
{
    initializer_t& v = get<L>();
    push_front(v);

    if constexpr (sizeof...(T) == 0)
    {
        if constexpr (std::is_invocable_r_v<X, L>)
        {
            v.reset = [x, lambda] { *x = lambda(); };
            return lambda();
        }
        else
        {
            v.reset = [x] { reset(*x); };
            return X{};
        }
    }
    else
    {
        v.reset = [x, t = std::make_tuple(val...)] { *x = std::make_from_tuple<X>(t); };
        return X(val...);
    }
}
```

**MY BIG RED BUTTON**



**IS BIGGER THAN YOURS**



# Thank You!

Nozomi Networks accelerates digital transformation by protecting the world's critical infrastructure, industrial and government organizations from cyber threats. Our solution delivers exceptional network and asset visibility, threat detection, and insights for OT and IoT environments. Customers rely on us to minimize risk and complexity while maximizing operational resilience.

[nozominetworks.com](http://nozominetworks.com)