

Rust for C++ developers

What this talk will NOT be about

- Rust syntax
- Enums
- Error handling
- Generics
- Macros
- Cargo
- Interoperability with C++

What this talk will be about

- What is memory safety
- How Rust ensures memory safety at compile time
- How to argue with the borrow checker

What is memory safety?

```
1  ✓ int main() {  
2      · char *buffer = new char[42];  
3      · delete [] buffer;  
4  
5      · // ...  
6  
7      · buffer[0] = 42; · // Boom!  
8      · return 0;  
9  }
```

What is memory safety?

```
1 ✓ int main() {  
2     char *buffer = new char[100];  
3     delete [] buffer;  
4  
5     // ...  
6  
7     buffer[0] = 42;  
8     return 0;  
9 }
```

```
1 int* g(int* a, int* b) { return b; }  
2  
3 int main()  
4 {  
5     int x;  
6     int* p = g(&x, nullptr);  
7     *p = 42; // boom  
8 }
```

What is memory safety?

```
1  ✓ int main() {  
2      char *buffer = new char[100];  
3      delete[] buffer;  
4  
5  
6  
7  
8  
9  }
```

```
1  int* g(int* p) {  
2      return p;  
3  }
```

```
4  std::string get_string()  
5  {  
6      return "A string literal casted to std::string";  
7  }
```

```
9  int main()  
10 {  
11     std::string_view sv = get_string();  
12     auto c = sv.at(0); // Boom  
13 }  
14
```

What is memory safety?

```
4  int main() {  
5      auto v = std::vector<int>();  
6      v.push_back(0);  
7      auto first_elem = &v[0];  
8      v.push_back(1);  
9      std::cout << "First elem is: " << *first_elem << std::endl;  
10     return 0;  
11 }
```

PROBLEMS OUTPUT TERMINAL DEBUG CONSOLE

```
> g++ .\push_back.cpp -o push_back.exe  
> .\push_back.exe  
First elem is: -2011292912  
> 
```

RAII

- Bind the life cycle of a resource to the lifetime of an object

```
std::mutex m;

void bad()
{
    m.lock();           // acquire the mutex
    f();                // if f() throws an exception, the mutex is never released
    if(!everything_ok()) return; // early return, the mutex is never released
    m.unlock();         // if bad() reaches this statement, the mutex is released
}

void good()
{
    std::lock_guard<std::mutex> lk(m); // RAII class: mutex acquisition is initialization
    f();                             // if f() throws an exception, the mutex is released
    if(!everything_ok()) return;     // early return, the mutex is released
}                                   // if good() returns normally, the mutex is released
```


RAII

- Example with memory

```
1
2
3 int main() {
4     auto vect = new int[30];
5
6     //....
7
8     delete vect;
9 }
10
```

```
1 #include <vector>
2
3 int main() {
4     auto vect = std::vector<int>(30);
5
6     //....
7
8
9 }
10
```

How Rust ensures safety

- Ownership
- Borrowing
- Lifetimes

Ownership

- Each value has an owner
- There can be **only one owner** at a time
- When the owner goes out of scope, the value will be dropped (released from memory)

The diagram shows a code snippet with three annotations. An arrow labeled 'Owner' points to the variable `s1` in line 2. An arrow labeled 'Value' points to the string literal `"Hello world"` in line 2. A third arrow points from the closing curly brace `}` in line 4 to the text 'End of scope s1 is dropped'.

```
1 fn main() {  
2     let s1 = String::from("Hello world");  
3     println!("{}", s1);  
4 }
```

Owner

Value

End of scope
s1 is dropped

Ownership

	C++	Rust
a = b	<ul style="list-style-type: none">- Depends on copy/move constructors- Mostly copies	<ul style="list-style-type: none">- If type is Copy, copy- Otherwise, move

- Assignment

- Primitive types or classes that implements Copy -> copy
- Otherwise -> transfers ownership (similar to std::move)

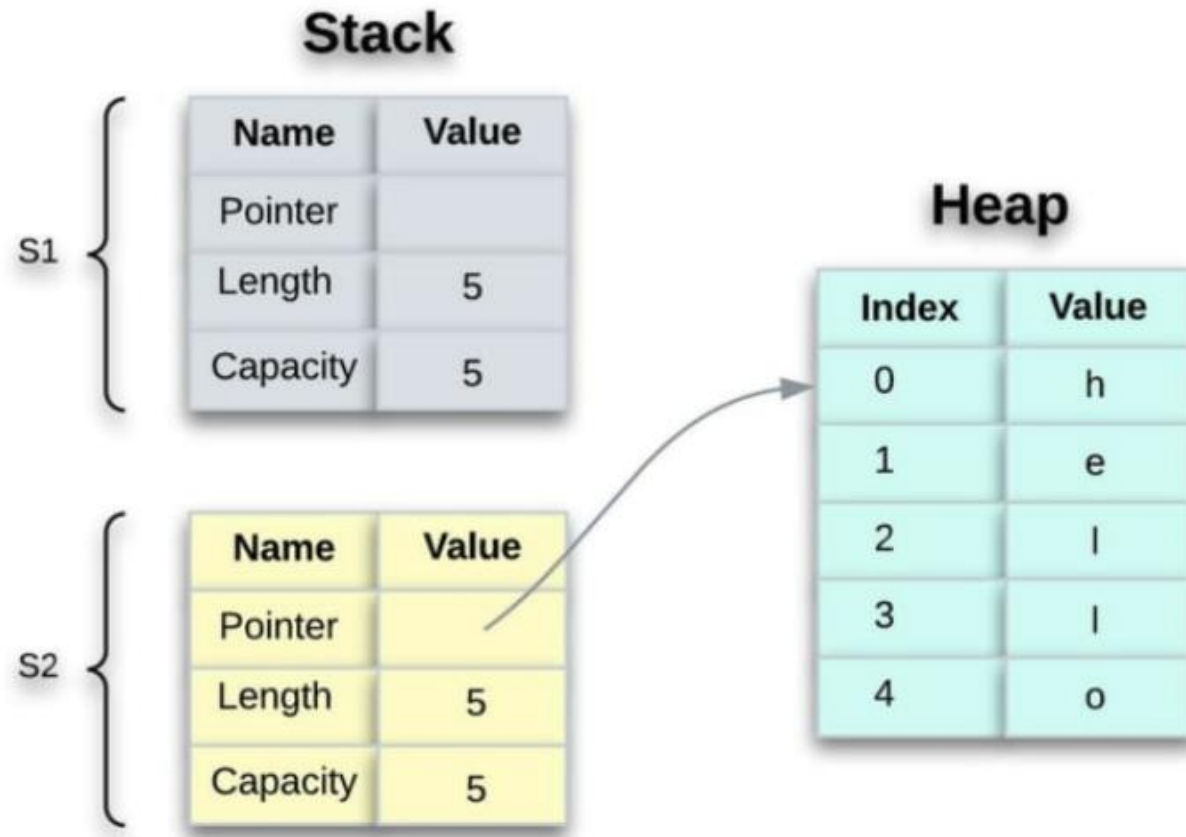
Example 1

```
let x = 5;
let y = x;
println!("{}", x, y); // OK -> primitive values are copied
```

Example 2

```
let s1: String = String::from("Hello world"); // 's1' is the owner of this string
let s2: String = s1; // Ownership is transferred --> s2 is the new owner and s1 is dropped
println!("{}", s1, s2); // Compiler error -> s1 is not valid anymore
```

Ownership



Ownership

Example 3

```
1  fn print_string(s: String) {  
2  |    println!("The string is: {}", s);  
3  }  
4  
   ► Run | Debug  
5  fn main() {  
6  |    let s: String = String::from("Hello world!");  
7  |    print_string(s);  
8  |  
9  |    println!("{}", s);  
10 }  
    }
```

Ownership

```
> cargo build
   Compiling ownership v0.1.0 (C:\Users\andre\Documents\Workspace\RustExperiments\ownership)
error[E0382]: borrow of moved value: `s`
  --> src\main.rs:9:20
6 |   let s = String::from("Hello world!");
  |   - move occurs because `s` has type `String`, which does not implement the `Copy` trait
7 |   print_string(s);
  |   - value moved here
8 |
9 |   println!("{}", s);
  |               ^ value borrowed here after move

= note: this error originates in the macro `$crate::format_args_nl` which comes from the expansion of the macro `println` (in Nightly builds, run with -Z macro-backtrace for more info)

For more information about this error, try `rustc --explain E0382`.
error: could not compile `ownership` due to previous error
> █
```

Borrowing

	C++	Rust
Mutable reference	<code>auto a = &b</code>	<code>let a = &mut b</code>
Immutable reference	<code>const auto a = &b</code>	<code>let a = &b</code>

- You can have immutable references

```
let s1: &String = &s;
```

- You can have mutable references

```
let s1: &mut String = &mut s;
```


- However, at any point in time, you can either have:

Many immutable
references

OR

1 mutable (exclusive)
reference

Borrowing



Borrowing

Example 4

```
let s: String = String::from("Hello world");

let s1: &String = &s;
let s2: &String = &s; // OK, multiple immutable borrows

println!("{}", s1, s2); // Hello world, Hello world
```

Example 5

```
let mut s: String = String::from("Hello world");

let s1: &mut String = &mut s; // OK, only one mut
*s1 = String::from("Bye bye"); // Can mutate string

println!("{}", s); // Bye bye
```

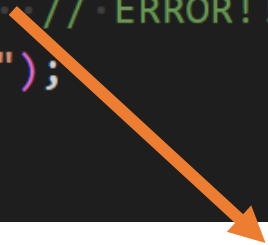
Borrowing

Example 6

```
let mut s: String = String::from("Hello world");

let s1: &String = &s;
let s2: &mut String = &mut s; // ERROR!!
*s2 = String::from("Bye world");

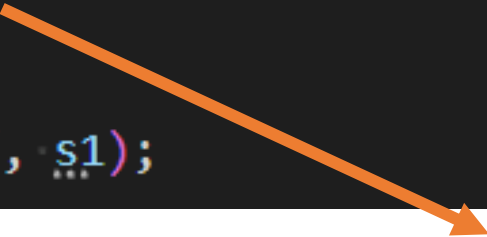
println!("{}", s1);
```



```
error[E0502]: cannot borrow `s` as mutable because it is also borrowed as immutable
--> src/main.rs:30:14
   |
29 |     let s1 = &s;
   |             -- immutable borrow occurs here
30 |     let s2 = &mut s; // ERROR!!
   |             ^^^^^^ mutable borrow occurs here
...
33 |     println!("{}", s1);
   |                   -- immutable borrow later used here
```

Lifetimes

```
let s1: &String;
{
    let s: String = String::from("Hello World");
    s1 = &s;
}
println!("{}", s1);
```



error[E0597]: `s` does not live long enough

--> src\lifetime_examples.rs:5:14

```
5 |         s1 = &s;
   |               ^^ borrowed value does not live long enough
6 |     }
   |     - `s` dropped here while still borrowed
7 |     println!("{}", s1);
   |                   -- borrow later used here
```

For more information about this error, try `rustc --explain E0597`.

error: could not compile `ownership` due to previous error

Lifetimes

- Lifetime = how long a reference lives (its scope)
- The lifetime of the borrower cannot outlive the borrowed value's owner

```
1 fn test_lifetime_1() {  
2     ... let s1: &String;  
3     ... {  
4         ... let s: String = String::from("Hello World");  
5         ...  
6         ... s1 = &s;  
7         ... }  
8         ...  
9         println!("{}", s1);  
10    }
```

Owner

Reference

Going back to safety

```
1  ✓ int main() {  
2      char *buffer = new char[42];  
3      delete [] buffer;  
4  
5      // ...  
6  
7      buffer[0] = 42; // Boom!  
8      return 0;  
9  }
```

```
let mut v: Vec<char> = Vec::with_capacity(42);  
drop(v);  
  
// ...  
v[0] = 'c';  
println!("{}", v[0]);
```

error[E0382]: borrow of moved value: `v`

--> src\cpp_examples.rs:7:5

```
3 | let mut v: Vec<char> = Vec::with_capacity(42);  
  | ----- move occurs because `v` has type `Vec<char>`, which does not implement the `Copy` trait  
4 | drop(v);  
  |     - value moved here  
...  
7 | v[0] = 'c';  
  | ^ value borrowed here after move
```

Going back to safety

```
1  int* g(int* a, int* b) { return b; }
2
3  int main()
4  {
5      int x;
6      int* p = g(&x, nullptr);
7      *p = 42; // boom
8  }
```

- No "nullptr" in Rust
- Would fail because of lifetime check
- "p does not live long enough"

Going back to safety

```
auto v = std::vector<int>();  
v.push_back(0);  
auto first_elem = &v[0];  
v.push_back(1);  
std::cout << *first_elem << std::endl;
```

```
let mut v: Vec<i32> = vec![];
v.push(0);
let first_elem: &i32 = &v[0];
v.push(1);
println!("{}", *first_elem);
```

```
error[E0502]: cannot borrow `v` as mutable because it is also borrowed as immutable
--> src\cpp_examples.rs:43:5
```

```
42 | let first_elem = &v[0];  
    | - immutable borrow occurs here
```

```
43 | v.push(1);  
    | ^^^^^^^^^ mutable borrow occurs here
```

```
44 | println!("{}", *first_elem);
    | ----- immutable borrow later used here
```


Smart pointers

	C++	Rust
Unique reference	<code>unique_ptr<T></code>	<code>Box<T></code>
Shared reference	<code>shared_ptr<T></code>	<code>Rc<T></code>

```
let s: MyStruct = MyStruct {x : 10};  
let a: Box<MyStruct> = Box::new(s); ... // Now a points to somewhere in the heap  
println!("{:?}", a);
```

```
let s: MyStruct = MyStruct {x : 10};  
let rc: Rc<MyStruct> = Rc::new(s);  
  
let s1: Rc<MyStruct> = rc.clone(); // Increments reference count  
let s2: Rc<MyStruct> = rc.clone(); // Increments reference count  
  
println!("{:?}", "{:?}", s1, s2);
```

Runtime borrow checking

- `RefCell<T>` is a smart pointer that checks borrows at runtime
- It can give mutable reference from immutable object (!!)
- This is known as "Internal Mutability Pattern"

```
let x: RefCell<MyStruct> = RefCell::new(MyStruct {x: 10});  
  
// Returns &MyStruct. Panics if value is currently mutably borrowed  
let x_ref: Ref<MyStruct> = x.borrow();  
  
// Return &mut MyStruct. Panics if value is currently immutably borrowed  
let x_ref_mut: RefMut<MyStruct> = x.borrow_mut();
```

Going multithreading

- A type is **Send** if it is safe to send it to another thread
- A type is **Sync** if it is safe to share between threads (T is Sync if and only if &T is Send)

	Single thread	Multiple threads
Sharing data	Rc<T>	Arc<T>
Allow mutability	RefCell<T>	Mutex<T>